JANNET AKANKSHA EKKA

AI Evaluation Engineer · LLM & Agent Testing · Red-Teaming · Test Automation

jannetekka96@gmail.com | +91 90788 02572 | linkedin.com/in/jannet-akanksha-ekka | github.com/JannetEkka | jannetekka.github.io/DSProjects | Kolkata, India · remote or relocation

PROFESSIONAL SUMMARY

I test AI systems. Four years leading QA automation at Deloitte taught me how software fails before the client sees it, and since 2024 I have applied that to LLMs and agents. When I graded the six agents in my own trading system on independent data, all six called direction at a coin flip, and I rebuilt the system around that result. I build the measurement before the feature: retrieval evals with significance tests, grounding checks that stop unsupported claims reaching a user, adversarial tests that judge an agent by the tool calls it made, and backtests against the obvious baseline. When the baseline wins, my write-up says so. Rank 1 in the UT Austin McCombs Post Graduate Program in AI/ML.

TECHNICAL SKILLS

LLM & Agent Evals	Labelled eval sets and harnesses, recall@k, MRR@k, macro AUC, paired bootstrap confidence intervals, held-out sets, baseline comparisons, LLM-as-judge, grounding checks on quotes, dates and amounts, counterfactual faithfulness tests, label-quality checks against expert labels
Red-Teaming	Tool-call interception with decoy tools, prompt injection written to read like ordinary work, “said no, did yes” detection, re-running approved fixes against the attacks that broke them, INCONCLUSIVE verdicts when nothing was classified
Statistical Validation	CPCV, Deflated Sharpe, PBO, FDR control, conformal prediction, forward tests on non-overlapping windows, replay of strategy variants through production code
Test Engineering	Test strategy, pytest (3,000+ tests on one system), Selenium, Katalon Studio, Jenkins and GitHub Actions CI, defect triage, liveness checks on production logs, review of AI-written code
AI Stack	Gemini (Vertex AI and AI Studio), Google ADK, Model Context Protocol, RAG (BM25, RRF, cross-encoders), FalkorDB, llama.cpp, Hugging Face, scikit-learn, PyTorch, TensorFlow, Claude Code
Languages	**Python and SQL are the languages I take into a coding interview.** TypeScript and JavaScript I have shipped and read fluently, and I would be ramping in them rather than leading

EVALUATION WORK

Smart Money Trading (SMT): an AI agent that has to prove itself · Founder & sole engineer

Jan 2024 – Present

Patent pending: Indian provisional application No. 202631090789, filed myself in July 2026

Live: smt-weex-trading-bot.jannet-ekka.workers.dev · *Public layer:* github.com/JannetEkka/smt-apac

- Graded my own multi-agent committee (six persona agents and a learned Judge) on independent market data. **All six called direction at a coin flip.** Rebuilt on that result: written rules with replayed records open every trade, and the agents became advisors graded on every call, earning a veto only after 30 proven warnings.
- Every explanation passes a **counterfactual test** before a user sees it: flip the advisor's input and re-run the decision; if the decision doesn't change, the reason is dropped.
- The learning loop retunes only through a **CPCV + Deflated Sharpe + PBO + FDR gate**. That gate still holds in shadow a forecaster that scored AUC 0.723 on a BTC forward test over 27 non-overlapping windows.
- Tested **1,800 strategy variants, 36 million simulated trades** from 2020 to 2026 at six exchanges' published fees, replayed through the production code. Nothing holding a position under four hours cleared its fees.
- Audited 98 AI-written merges line by line and found two failures the green test suite missed: an order-format bug that had **rejected 1,137 live orders over three months**, and a memory freeze that kept the daemon down 60% of the time for ten days. A fix now counts as done only when the production log shows it.
- A register of **611 measured findings** decides what ships. Any substituted default logs a warning and a liveness flag, a rule that has caught four multi-week data bugs. 91,000 lines of Python, 3,000+ tests.

Wiretrap: adversarial security testing for AI agents · Gen AI Academy APAC, Cohort 3

2026

Live: wiretrap.onrender.com · *Code:* github.com/JannetEkka/wiretrap

- Most scanners grade the agent's reply. Wiretrap reads the tool-call log. Every declared tool is cloned as a dead copy with the same name and JSON Schema that records the call and returns fake data; attacks run in separate Gemini conversations and a judge model reads the log.
- Caught a Google codelab agent saying “I don't have access or authorization to modify configuration files” after it had already run `ls -la` through its shell tool. Text scanners pass that; Wiretrap flags it with the arguments.
- A run with nothing classified reports INCONCLUSIVE, never HARDENED. Attacks on tools the agent doesn't have are dropped from the score. 8 attack categories, 9.4K lines of TypeScript.

SmartDesk: evaluated RAG · Gen AI Academy APAC hackathon, top 101 of 1,500+

2026

github.com/JannetEkka/smartdesk · *Eval write-up:* evals/RESULTS.md

- Built the eval harness before touching retrieval: 40 labelled questions over 120 notes, recall@k and MRR@k, paired bootstrap at 10,000 resamples. Baseline recall@5 **0.963**.
- A cross-encoder reranker won on the development embedder (MRR@10 +0.094, p=0.04), then fell to +0.028 with a CI spanning zero on the production embedder, at the cost of 190 MB of PyTorch and search going from 29 ms to 2,382 ms. **The baseline shipped.**

Klarbrief and Chithi: grounding checks on LLM explanations · Hack Apertus; Meta AI Developer Hackathon

Oct – Nov 2026

Code: github.com/JannetEkka/klarbrief · *Chithi:* chithi-oprj.onrender.com

- Apertus 1.5 8B explains official Swiss letters in the reader's language, offline on a 4-core CPU. Before any point is shown, a deterministic check finds its quoted words in the letter and verifies its date or CHF amount.
- Measured on 48 explanations of 28 synthetic letters plus 8 held-out letters: 48/48 in the reader's language, 28 of 29 real deadlines shown, **0 wrong CHF amounts shown**; the check stopped 6 of the model's 7 wrong dates. No wrong date or amount shown on the held-out set. 34 tests, 5-page report.
- **Chithi** (Meta hackathon, SAM 3.1 + Muse Spark): photos of Indian letters explained in Bengali, Hindi or English, each value checked against the quote and a second reading. **0 of 114** shown wrong on 28 live runs (16 hard photos); 53/53 planted errors held back; a key-facts metric added after a demo run left out the deadline.

Lore: backtesting an agent against the obvious baseline · Graph Hacks, FalkorDB × WeMakeDevs

Oct 2026

Live: lore-gh5y.onrender.com · *Code:* github.com/JannetEkka/lore

- A Gemini agent with 12 Cypher tools answers “who should review this?” from a FalkorDB graph of an open-source project's history, citing the PRs behind every answer. On 300 real PRs it had a real reviewer in its top three **90%** of the time, against 88% for “the most frequent reviewer”. On the single top pick the baseline wins, and the write-up says so.

More evaluation in October 2026 · Kaggle RSNA Knee · Monad Metropolis

Oct 2026

- **RSNA Knee Abnormality Detection** (Kaggle, 5,900+ teams, 12 labels from knee MRI): training labels had to be derived from free-text radiology reports, so I scored two candidate label sets against 58 radiologist labels (macro AUC 0.894 against 0.870) before retraining on the better one.
- **SMT Receipts**: my trading agent's calls are sealed on Monad testnet by an ownerless, verified contract and graded later against RedStone-signed prices net of fees, so its track record cannot be edited after the fact.

PROFESSIONAL EXPERIENCE

Deloitte Consulting LLP · Test Automation Engineer (Analyst) *Aug 2019 – Jan 2023 · Bangalore*

- **Led a 6-member QA engineering team** on automated testing for Fortune 500 clients including AT&T Consumer Connect, Hewlett Packard Enterprise and The Dental Supply Company.
- Built Selenium and Katalon Studio automation that **improved test execution efficiency by 83%**; Jenkins-based reporting cut manual analysis time by **75%**.
- Wrote Python analysis and Tableau dashboards tracking quality across **343 business components**, cutting test failures by 25%; analysed **50,000+ SAP Hybris transactions** to isolate performance bottlenecks.
- Client-facing throughout: ran demos, working sessions and requirement calls with client teams in US, UK and India hours, and covered weekend on-call.
- Joined as a QA automation intern in Jan 2019 (North Dakota Health and Human Services website, POEDI XML testing); converted to analyst after six months. 2 Spot Awards and an Applause Award.

Tech: Selenium, Katalon Studio, Python, JavaScript, Jenkins, Salesforce, SAP Hybris, Tableau, JIRA, Git

JTech Solutions (pre-incorporation) · Founder & AI Engineer *Jan 2024 – Present · Kolkata*

- Building SMT and SMT World (above) as the company's first product, while a full-time family carer and alongside the PGP in AI/ML.

AutoKorrekt · Lead Frontend Developer (EdTech startup MVP) *May 2023 – Jan 2024*

- Led the frontend of an AI answer-evaluation platform for teachers: OCR over **AWS Textract**, AI evaluation and validation per student, English/German feedback on Amazon Comprehend; scaled to **1,000+ concurrent submissions**.

EDUCATION

Post Graduate Program in AI & Machine Learning · UT Austin McCombs & Great Lakes *Mar 2025*

Rank 1 in batch · GPA 4.09/5

Bachelor of Technology, Information Technology · KIIT University, Bhubaneswar *2015 – 2019*

CGPA 7.17/10

ACHIEVEMENTS & CERTIFICATIONS

- **Provisional patent filed**, India, App. No. 202631090789 (Jul 2026), covering the SMT decision and verification architecture.
- **WEEX AI Wars I: #1 on the preliminary leaderboard** of 230+ teams, then one of 37 finalists in the live AI trading competition; featured in WEEX's coverage.
- **Gen AI Exchange Hackathon 2025**, problem statement “Automating Test Case Generation with AI”.
- **Winner, Best DeFi Application**, OpenServ × Hack2skill hackathon; runner-up, Best Website Application.
- **Top 101 of 1,500+ teams**, Google Cloud Gen AI Academy APAC 2026 Cohort 1 hackathon (SmartDesk).
- **CPSAT**, Certified Professional Selenium Automation Testing, 2022. Google Cloud Gen AI Academy 2025 and APAC 2026 Cohorts 1–3. All on Credly and Google Skills.